

DGTP SINT MAARTEN • NON-CRM INTEGRATION PROGRAMME

VROMI & Kadaster Integration Readiness

Address Registry & Land Registry

Stakeholder Briefing

40-Minute Presentation Script

Enterprise Architect – Orange BD

7 May 2026 • v1.0 • Confidential

Presentation Overview

Slide	Section	Duration
1	Title - Programme Context & Introductions	3 min
2	Executive Summary - 37 items, 0 confirmed	5 min
3	Impact Playground - Live Blocker Demonstration	7 min
4	Master Integration Flow - Both paths blocked	5 min
5	API & Authentication - Token chain failure	5 min
6	Data Schema & CMF - Field mapping blocked	4 min
7	Sandbox & Pagination - Sprint start blocked	4 min
8	Legal & Governance - Article 22 critical path	4 min
9	By Audience - Ownership & Action Plan	3 min
	Q&A Buffer	—
Total		40 min

Before you begin: Confirm the dashboard is open on the projector, logged in (username: jahidul, password: orange), and showing Slide 1. Have the AI chat panel closed. Ensure all stakeholders are in the room – you need VROMI IT, Kadaster IT, DLT/NRPB, GoSM ICT, GoSM Legal, and Orange BD represented.

SLIDE 1 — Programme Context & Introductions

[0:00 -- 3:00]

SLIDE 1 – Title: VROMI & Kadaster Integration Readiness

[Slide is visible. Walk to the front. Let the room settle. Speak calmly and with authority.]

Good morning, everyone. Thank you for making time today.

My name is Jahidul, and I am the Enterprise Architect on the DGTP Sint Maarten programme – working with Orange BD on the technical integration between the government’s registries and the new citizen portal platform.

What I am presenting today is not a status report in the traditional sense. It is a readiness assessment. Specifically, it is an honest, structured view of where we stand right now in our ability to integrate the VROMI Address Registry and the Kadaster Land Registry into our platform – and what is blocking us from starting.

I want to set the tone clearly: **this presentation contains no good news**. Every single item on our integration checklist is still open. Zero confirmed. And I believe every person in this room has a role in changing that.

By the time we leave this meeting, I want three things to happen:

1. Every checklist item has a **named owner** – a specific person, not a team.
2. The two longest-lead items – legal review and network access – have been **formally initiated**.
3. Orange BD leaves with confirmed **endpoint URLs** from both VROMI IT and Kadaster IT.

[PAUSE] Everything I present today is linked to the Non-CRM Integration Readiness Checklist that was shared earlier. This dashboard visualises all 37 items in that document. Let us begin.

[ADVANCE SLIDE]

SLIDE 2 — Executive Summary

[3:00 -- 8:00]

SLIDE 2 – Executive Summary: 37 items, 0 confirmed

[Gesture to the three stat chips at the top: 24 critical, 9 high risk, 2-4 weeks legal.]

This slide is the most important slide in the deck. I want you to look at these three numbers.

Twenty-four critical blockers. Nine high-risk items. And a two-to-four-week lead time on legal review that has not yet been started.

To be clear about what “critical blocker” means: it means that without resolving this item, the integration adapter *cannot be built at all*. Not built slowly. Not built with risk. **Cannot be built.** A critical blocker is an architectural dead end.

Let me walk you through the sprint timeline you see on the left side of this slide.

Right now, today – there is exactly one thing Orange BD can do without waiting: build the APIM mock policy. That is the only unblocked track. Everything else requires inputs from people in this room.

Week one – if nothing changes, the adapter build remains completely blocked. We have no URL to connect to. We have no authentication mechanism to implement. We cannot write a single line of adapter code against a live system.

Week two – the sandbox is still unavailable. We cannot test safely. Our only option would be to connect directly to production registries, which violates data protection rules and creates a real risk of corrupting live government records.

Week three – Article 22 legal review is still running, assuming it started today. If it has not started today, we are now in week five or six before it concludes.

Week four – Sprint 1 start date is missed. The programme timeline shifts by a minimum of four weeks.

[PAUSE] Now look at the right panel – the stakeholder action rows. I want each of you to find your organisation and read what is listed next to it. These are not suggestions. These are the **minimum inputs required** for your part of the integration to proceed.

Pause here and give the room 20-30 seconds to read the stakeholder rows. Make eye contact with the lead from each organisation as they read. Let the silence do the work.

The good news – and there is a small piece of good news – is that resolving the top three items alone would move us from zero percent readiness to over sixty percent. Those three items are: §1.1 the VROMI base URL, §1.4 the Kadaster base URL, and initiating the Article 22 legal review. Those three actions, if taken today, change the trajectory of this programme immediately.

[ADVANCE SLIDE]

SLIDE 3 — Impact Playground

[8:00 -- 15:00]

SLIDE 3 – Impact Playground: What breaks if this stays unresolved?

[This is the most interactive slide. All toggles are ON at the start. Walk through a live demonstration.]

I want to do something a little different for this slide. Rather than telling you what is broken, I am going to show you – live – what happens when specific blockers remain unresolved. And then I am going to show you what happens when we resolve them.

You can see on the left a list of eight blockers, each with a toggle switch. On the right, a live impact canvas shows the cascading downstream failures for each active blocker. At the bottom is a sprint readiness score.

Right now, with the six critical blockers all active, our readiness score is **five percent**. Let me demonstrate what that means.

Live demo sequence - perform slowly, narrate each action:

1. Point to the flow diagram showing the URL blocker - Platform to Adapter, dashed red line, blocked.
2. Click the "Endpoint URL missing" toggle OFF.
3. Watch the score jump. Say: "Resolving just this one item moves us from 5% to 20% - because it unblocks eight downstream dependencies simultaneously."
4. Click "Auth method unknown" toggle OFF.
5. Say: "Authentication resolved - now at 35%."
6. Toggle ALL critical items OFF one by one.
7. Score reaches 100%. Say: "This is what full readiness looks like."
8. Toggle everything back ON. Score drops to 5%.

What you have just seen is not a theoretical model. Every node in that flow diagram – every dashed red edge – represents a real architectural dependency in our integration design. When the URL is missing, the adapter has **no target**. It is like building a road without knowing where it needs to go.

Let me draw your attention to the flow diagram tabs at the top of the canvas. I can switch between different blocker flows. Let me show the authentication chain.

Click the “Authentication” tab. Point to the token endpoint node. Narrate: “Every single API call to ArcGIS requires a valid token. The token comes from an endpoint called generateToken. We do not have this URL. Without it, every call returns 401 Unauthorised. The adapter cannot authenticate against either registry.”

Now let me show the legal flow.

Click the “Legal clearance” tab. Point to the Article 22 node. Say: “This is the longest lead-time item in the entire programme. The review takes two to four weeks minimum – and it cannot be expedited. The review must be initiated today for us to have any chance of Sprint 1 starting on schedule.”

I want to leave you with one key insight from this slide: **the order in which we resolve blockers matters enormously**. Resolving the endpoint URL first is the highest-leverage action available to us right now. It has eight downstream dependencies. No other single action has that multiplier effect.

[PAUSE] You will also notice the What-If AI bar at the top of this slide – it automatically narrates the impact analysis as I toggle blockers. This is a feature of the dashboard for your team to continue exploring after today.

[ADVANCE SLIDE]

SLIDE 4 — Master Integration Flow

[15:00 -- 20:00]

SLIDE 4 – Master Integration Flow: All paths blocked

[Gesture to the architecture diagram showing APIM, two adapters, ArcGIS registries, and Cosmos DB.]

This slide shows our target integration architecture and the current state of every connection within it.

Let me describe what you are looking at. On the left: our integration platform, mediated by APIM – Azure API Management. In the middle: two adapter services – one for VROMI Address Registry, one for Kadaster Land Registry. On the right: two ArcGIS Feature Service instances. And at the far right: Cosmos DB, our canonical data store.

Every arrow between these components should be a solid green line. To-day, every arrow is a **dashed red line**. Every connection is unconfirmed.

The VROMI adapter cannot be built because we do not have the ArcGIS Feature Service URL. The adapter needs to know where to send its queries. Without the URL, there is nothing to build against.

The same is true for Kadaster. §1.3 and §1.4 in the checklist cover this. Both are Open. Both have been Open since the checklist was created.

I want to specifically address the change notification strategy, shown in the third section of this slide. We do not know whether either ArcGIS instance supports webhooks or the changeTracking capability. This matters because it determines whether our platform can receive *push notifications* when data changes, or whether it must *poll* the registries continuously.

Polling is expensive. It creates constant load on the ArcGIS servers. It produces data that is always slightly stale. And it requires an agreed polling interval with each registry team – which we also do not have.

The ideal answer is webhooks with changeTracking enabled. The fall-back is polling. We cannot design either solution until §5.1 through §5.3 are confirmed.

Critical path reminder: The single most impactful action from this slide is VROMI IT and Kadaster IT confirming their ArcGIS Feature Service base URLs. This needs to happen before we leave this meeting, or assigned as a 48-hour action item with a named owner and a due date.

[ADVANCE SLIDE]

SLIDE 5 — API & Authentication

[20:00 -- 25:00]

SLIDE 5 – API & Authentication: Token, accounts, credentials all missing

[Reference the stat chips: 5 auth items open, 0 credentials provisioned, 0 TLS confirmed.]

Authentication is the second most critical section in the checklist. Even if we resolve the URL question today, we immediately hit the authentication wall.

ArcGIS uses token-based authentication. The mechanism works like this: the adapter sends a POST request to a generateToken endpoint with service account credentials. ArcGIS returns a token. That token expires – by default after sixty minutes. The adapter must then refresh the token automatically before it expires, or every API call starts returning 401 Unauthorised errors.

This is well-understood ArcGIS architecture. **But we cannot implement any of it** because we do not know the URL of the generateToken endpoint. That URL is §12.2 in the checklist. It is Open.

Beyond the token endpoint, there are three related blockers I want to highlight.

First - service accounts. Item §2.5. Both registries must provision dedicated service accounts for integration purposes. Using personal user accounts is not acceptable. It violates security policy, creates audit trail gaps, and breaks immediately when a staff member leaves or changes roles. We need named service accounts. Neither has been provisioned.

Second - credentials delivery. Item §2.6. Once service accounts are created, the credentials must be delivered securely. Plain email is not acceptable. We need an agreed secure channel – a password manager like KeePass or Bitwarden. This has not been agreed.

Third - TLS version. Item §2.7. Our integration platform enforces TLS 1.2 as a minimum. If either ArcGIS instance still supports TLS 1.0 or 1.1, all connection attempts will be rejected at the handshake level before a single API call is made. We need confirmation from GoSM ICT that both instances meet this standard.

Pause here. Turn to the GoSM ICT representative specifically. Say: “I need TLS confirmation from your team. This is a quick check – it takes minutes once someone looks at the ArcGIS server configuration. Can you commit to providing that answer within 24 hours?”

And finally – IP allowlisting. §2.8. Our integration platform has specific outbound IP addresses. Those IPs must be whitelisted in the government firewall. The process for requesting this typically takes one to three weeks. We need to submit that request today. And importantly – we need *two* requests: one for the development environment, one for production. These are different IP ranges.

[ADVANCE SLIDE]

SLIDE 6 — Data Schema & CMF

[25:00 -- 29:00]

SLIDE 6 – Data Schema & CMF: Field mapping blocked

[Reference the field mapping diagram: ArcGIS fields on the left, CMF fields on the right, question marks in the middle.]

The Canonical Message Format – the CMF – is the core data contract for our integration. It defines exactly what fields our platform stores, in what format, and with what constraints. The CMF cannot be fully designed until we know the ArcGIS field names.

Let me focus on two fields that are particularly critical.

The **A-key** – item §12.4. This is the primary key for address records in the VROMI ArcGIS system. In our CMF, it maps to the `addressId` field in `AddressRecord A1`. Every other registry that references an address – `PersonRecord`, `BusinessRecord` – links to this identifier.

The problem: we do not know what the ArcGIS column is called. It might be `OBJECTID`. It might be `ADDRESS_ID`. It might be `A_KEY`. We do not know. And this is not a minor detail – if we build our CMF against the wrong field name, every cross-registry lookup in the entire platform fails silently. No errors. Just missing data.

The **PR1 field** – item §12.7. This is the permit number field in the Kadaster ArcGIS layer. Same problem. Unknown name. But there is an additional risk here: the PR1 field must be *stable across ArcGIS data refreshes*. If Kadaster's ArcGIS database is refreshed periodically and the permit IDs change or are renumbered, all existing permit records in Cosmos DB become unreachable.

Turn to the Kadaster IT representative. Ask directly: “When Kadaster refreshes its ArcGIS data, do permit record IDs remain stable? Is the permit number a natural key or a synthetic key? This answer affects our entire data persistence strategy.”

There are two additional technical concerns on this slide I want to flag quickly.

Coordinate reference system – §12.6. Our platform requires WGS 84, EPSG:4326. If VROMI's ArcGIS uses a local Caribbean CRS, we need a projection transformation in the adapter. This is not complex, but it adds development time and a library dependency. We can confirm this the moment we have access to the schema endpoint.

Cross-registry parcel ID linkage – §12.8. Land parcel records from

Kadaster must link to business records in COCI via a shared parcel identifier. We have not confirmed that Kadaster and COCI use the same parcel ID format. If they differ, cross-registry queries produce silent mismatches.

Both of these can be resolved with one action: granting Orange BD read access to /FeatureServer/0?f=json on both ArcGIS instances.

[ADVANCE SLIDE]

SLIDE 7 — Sandbox & Pagination

[29:00 -- 33:00]

SLIDE 7 – Sandbox & Pagination: Sprint start blocked

[Reference the stat chips: 3 critical blockers, 0 sandboxes confirmed, 1 can proceed.]

Without a sandbox environment, we cannot start the development sprint safely. This is not a preference. It is a constraint imposed by data protection obligations.

If our adapters connect to production registries during development, we risk making queries that could expose or modify live government records. Sint Maarten's land and address data is legally sensitive. A development error that corrupts production data in Kadaster is not recoverable quickly. The impact on government operations would be significant.

We need staging environments. Specifically:

For **VROMI**: a staging ArcGIS endpoint with at least fifty representative test address records. The test data should include Sint Maarten-style addresses with postcodes, coordinate data, and realistic field values.

For **Kadaster**: a staging endpoint with at least three parcel records that are linked to permit records. This is the minimum we need to verify the PR1 linkage and cross-registry join logic.

Item §6.2 covers VROMI. Item §6.3 covers Kadaster. Both are Open. Neither has a named provisioning owner.

Now let me address the **pagination** issue, because this one is subtle and dangerous.

ArcGIS Feature Services have a built-in query result limit called `maxRecordCount`. The default is typically between one thousand and two thousand records, depending on the ArcGIS version. If you query a layer with ten thousand records and do not implement pagination, ArcGIS returns the first one thousand records – **and returns no error**. The response looks complete. The data is not.

This is a silent failure mode. Our platform would store one thousand addresses and believe it has the complete dataset. Reports generated from that data would be subtly wrong – wrong counts, missing records, incorrect coverage calculations.

The fix is straightforward: we implement pagination using `resultOffset` and `resultRecordCount` parameters. But we need to know the actual `maxRecordCount` setting on each ArcGIS instance before we can design this correctly.

The one thing that can proceed: APIM mock policy. Orange BD can build the complete APIM mock contract using the expected GeoJSON schema today, without waiting for any registry inputs. This unblocks the citizen portal frontend development immediately. We are doing this regardless.

[ADVANCE SLIDE]

SLIDE 8 — Legal & Governance

[33:00 -- 37:00]

SLIDE 8 – Legal & Governance: Article 22 must start today

[Speak slowly and deliberately. This is the most time-sensitive section. Look directly at the DLT/NRPB and GoSM Legal representatives.]

I want to spend a few minutes on legal and governance, because this section contains the **critical path items for the entire programme timeline**. Not the most technically interesting items – the most time-critical ones.

Sint Maarten's Electronic Administration Ordinance – the EA Ordinance – Article 22 governs cross-registry data sharing. It defines which sub-registries can share data, under what conditions, and whether citizen consent is required for specific combinations.

The Article 22 review must be completed before our integration platform can legally receive data from VROMI or Kadaster. This review **takes two to four weeks minimum**. It cannot be expedited. The clock starts on the day the review request is formally submitted to GoSM Legal.

The review has not been started. Item **§9.4**. Still Open.

If we initiate the review today, we have a reasonable chance of legal clearance by week four or five, which aligns with our sprint timeline. If we leave this meeting without initiating the review, Sprint 1 is delayed by the same number of days we wait.

Critical path statement: Every day Article 22 review does not start is a day directly added to the Sprint 1 start date. This is not a risk. It is a certainty. DLT/NRPB must contact GoSM Legal today.

Beyond Article 22, there are three governance items I want to flag:

Restricted fields list - §9.2. GoSM Legal must provide a list of fields in each registry that may not be read by the integration platform. This could include national ID numbers, sensitive personal data, or legally protected information. Without this list, our CMF field mapping may inadvertently include restricted fields – creating a legal liability and requiring an architectural redesign.

Master data ownership - §9.1. Both the Civil Registry and the Address Registry contain address data. If they conflict, our platform needs to know which source is canonical. This must be decided at a governance level, not a technical level. The decision affects our Cosmos DB write logic and conflict resolution strategy.

Cache TTL - §9.5. Personal data cache duration must be defined before Cosmos DB can be configured. Some personal data may require zero-cache – meaning every query must hit the live registry. Others may be cacheable for hours or days. GoSM Legal and DLT/NRPB need to define this policy jointly.

[ADVANCE SLIDE]

SLIDE 9 — By Audience: Ownership & Action Plan

[37:00 --

40:00]

SLIDE 9 – By Audience: 37 items, 0 named owners

[This is the closing slide. Shift tone from analytical to action-oriented. Stand straighter.]

This is the final slide, and it is the most important one for what happens after we leave this room.

You can see on this slide the stakeholder matrix – every organisation represented here today, with the checklist items that are specifically assigned to them.

The critical observation: **thirty-seven items are Open, and not a single one has a named owner.** “VROMI IT” is not an owner. “GoSM Legal” is not an owner. An owner is a specific human being – a name, a contact, and a delivery date.

Before we close, I am asking each team lead to do one thing: **identify the named person responsible for their top three items and commit to a delivery date.**

Let me summarise the five most time-critical actions:

1. **VROMI IT and Kadaster IT - within 48 hours:**

Provide the ArcGIS Feature Service base URLs. This is the single highest-leverage action available. Eight downstream items unblock immediately.

2. **DLT/NRPB - today, before close of business:**

Contact GoSM Legal to formally initiate the Article 22 review. Name the person who will own this action.

3. **GoSM ICT - this week:**

Submit IP allowlist requests for both dev and production environments. Confirm whether ArcGIS endpoints are internet-facing or VPN-only.

4. **GoSM Legal - within 2 weeks:**

Provide the restricted fields list. Begin Article 22 review upon DLT/NRPB's formal request.

5. **Orange BD - starting now:**

Build APIM mock policy and Postman collection. Begin CMF field mapping draft with TBC placeholders for unknown fields.

I want to close with a direct statement.

The integration architecture is sound. The technical design is ready. Orange BD has the capability and the capacity to build this. **The only thing blocking us is information that exists in this room right now** - URLs, field names, legal decisions, network configurations.

Every day those items remain unresolved, the programme slips. Not because of engineering delays. Not because of platform limitations. Because the registries have not yet provided what we need to begin.

I am asking for your commitment today. Not next week. Today.

Thank you.

[PAUSE] *[Open the floor to questions. You have approximately 10-15 minutes for Q&A. Key anticipated questions are listed on the following page.]*

Anticipated Questions & Prepared Answers

These are the questions most likely to arise during Q&A. Prepare brief, confident answers. Do not wing these – they will be asked.

Q: Why can we not just connect to production while the sandbox is being set up?

Because Sint Maarten government data is legally protected and operationally critical. A development error that queries or modifies production address or land records could create incorrect entries that affect legal transactions – property sales, permit applications, business registrations. The sandbox is not a preference; it is a data protection obligation. We will not connect to production for development under any circumstances.

Q: How long will the integration actually take once we have the inputs?

Once we have confirmed URLs, authentication details, sandbox access, and CMF field names, the core adapter development is approximately three to four weeks. APIM integration, testing, and hardening adds another two weeks. Total: five to six weeks from input confirmation to a production-ready adapter. The legal review is the critical path – it runs in parallel but takes two to four weeks regardless of development speed.

Q: Can we integrate with just one registry first and come back to the other?

Technically yes, but the citizen portal's core queries join address data with land data. A partial integration produces a portal that cannot answer the questions citizens will actually ask. We recommend proceeding with both in parallel – the adapters are architecturally independent, so there is no cost to building them simultaneously.

Q: Is the ArcGIS data accurate enough for the portal?

Data quality is outside the scope of today's presentation – that is a conversation for the registry teams. Our adapter will store what it receives. If there are known data quality issues in the source registries, those should be flagged separately and addressed before go-live. The CMF does include validation logic that will flag missing required fields during ingestion.

Q: What happens if Article 22 is not cleared in time?

The integration platform cannot legally receive data from the registries until Article 22 is cleared. Sprint 1 development can proceed against mock data, but no real registry data can flow until legal clearance is obtained. If this takes longer than four weeks, we reassess the programme timeline. There is no workaround.

Q: Why is the dashboard interactive – is this a production tool?

The dashboard is a readiness visualisation built specifically for this programme. It is not a production monitoring tool. It serves as a single source of truth for the checklist status, a communication tool for stakeholder meetings like this one, and a reference document for the integration design. It will be maintained and updated as items are resolved.

Quick Reference – Key Section Numbers

Infrastructure

- §1.1 VROMI base URL
- §1.3 Excel to ArcGIS migration status
- §1.4 Kadaster layer IDs
- §1.5 Network access type (VPN?)

Authentication

- §2.1 Auth mechanism
- §2.5 Service account provisioning
- §2.6 Credential sharing channel
- §2.7 TLS version confirmation
- §2.8 IP allowlist
- §12.2 Token endpoint URL

Data Schema

- §3.2 VROMI data dictionary
- §3.4 Kadaster layer schemas
- §12.4 A-key field name
- §12.6 Coordinate reference system
- §12.7 PR1 permit field name
- §12.8 Parcel ID COCI linkage

Sandbox & Testing

- §6.1 Sandbox requirement
- §6.2 VROMI staging endpoint
- §6.3 Kadaster test parcels
- §6.4 Credential provisioning owner

Change Notification

- §5.1 Webhook support
- §5.2 changeTracking enabled
- §5.3 Polling interval
- §8.2 maxRecordCount

Legal & Governance

- §9.1 Master data ownership
- §9.2 Restricted fields list
- §9.3 Read-only confirmation
- §9.4 Article 22 legal review
- §9.5 Cache TTL policy
- §9.6 Data classification

Engagement Slide Notes — Facilitator Reference

The seven engagement slides appear automatically in the deck after each main slide. They are styled differently (orange gradient header, “QUICK CHECK-IN” eyebrow) so the audience recognises a shift in mode. Each has 3–4 questions and a facilitator guidance box visible only to you on this sheet. Three response buttons at the bottom let you mark: “Room is aligned”, “Flag for follow-up”, or “Action assigned”.

Engagement 1 — After Executive Summary

“Before we go further — a few things to confirm”

- Does this picture match your understanding of where the integration stands?
- Is there a champion in your organisation who can own delivery of your items?
- Has the Article 22 review been discussed internally before today? **DLT/NRPB**
- Are there constraints on the registries not in the checklist?

Guidance: Let the room respond freely. Listen for surprises and identify which teams are engaged versus passive. Do not rush this — the best information in the meeting often surfaces here.

Engagement 2 — After Impact Playground

“You have seen the cascading impact — what resonates?”

- Does the dependency chain for the URL blocker reflect how your ArcGIS is actually structured?
- Does Kadaster’s ArcGIS use changeTracking, or is polling the only option? **§5.1**
- What would it take from your team to move the readiness score past 60%?

Guidance: Invite VROMI and Kadaster IT leads to walk up to the screen. Make them contributors to the solution, not an audience.

Engagement 3 — After Master Integration Flow

“The architecture is ready — we just need the connection details”

- Is the VROMI ArcGIS migration from Excel fully complete, or is data still split? §1.3
- Can you share the base URL today — even a rough one? §1.1 §1.4
- GoSM ICT — is ArcGIS accessible from outside the government network? §12.10

Guidance: Direct the URL question specifically at named VROMI IT and Kadaster IT representatives. Pause. Let them think. If they need to check, ask for end-of-day today and write the commitment down.

Engagement 4 — After API & Authentication

“Authentication is the second wall — let’s see where you are”

- Is ArcGIS Portal configured, or vanilla ArcGIS Server with token auth? §2.1
- Can a dedicated service account be provisioned this week? §2.5
- GoSM ICT — do you have visibility on the TLS version running on ArcGIS? §2.7

Guidance: Frame service account provisioning as “a non-personal login with read-only access” — not a complex security task. The goal is yes or no, not a technical debate.

Engagement 5 — After Data Schema & CMF

“The field names are the key — literally”

- Does VROMI know the primary address identifier field name in ArcGIS? §12.4
- When Kadaster refreshes ArcGIS data, do permit IDs remain stable? §12.7
- Does anyone know if coordinate data is in standard GPS format? §12.6

Guidance: Show a coordinate pair example — 18.0, -63.0 (WGS84) versus 442000, 2000000 (local). Makes the CRS question immediately graspable for non-technical stakeholders.

Engagement 6 — After Sandbox & Pagination

“The sandbox is our safety net — let’s agree on what it needs to contain”

- Does a staging ArcGIS environment already exist, even informally? §6.1 §6.2

- Kadaster — can you export a small set of test parcels with linked permits? §6.3
- Does anyone know the maxRecordCount on the ArcGIS services? §8.2

Guidance: Frame the sandbox as protecting *their* live data, not as a developer convenience. Both VROMI and Kadaster have strong incentives to prevent production corruption.

Engagement 7 — After Legal & Governance

“Legal clarity protects everyone — including the registries”

- DLT/NRPB — who is the named person initiating Article 22 review today? §9.4
- GoSM Legal — are there obviously restricted fields we should be designing around? §9.2
- Which is the canonical address master source: Civil Registry or Address Registry? §9.1

Guidance: The Article 22 naming question is the most important ask in the presentation. If DLT/NRPB cannot name a person, escalate: “Who in this room has the authority to assign this right now?”

Closing Engagement — After By Audience

“Before we close — commitments on the record”

- VROMI IT — ArcGIS base URL, or date by which you will provide it?
- Kadaster IT — URL and PR1 field name. Your delivery date?
- DLT/NRPB — Article 22. Named person. Today. Confirmed?
- GoSM ICT — IP allowlist and TLS confirmation within 24 hours?

Guidance: Go around the table. Ask each lead to state their commitment out loud. Write it down visibly. This is the most important 3 minutes of the meeting.

NEW: Scope Clarification Slide — Before Slide 1

SCOPE & CONTEXT SLIDE — Before the main presentation begins

[This slide appears first, before the Executive Summary. Dark id-bar. Left column: 6 registries. Right column: SVG architecture diagram.]

Before beginning the main presentation, spend **90 seconds** on the scope slide. Its purpose is to prevent the most common confusion: stakeholders assuming this session covers all six registries, or that it covers the functional CRM.

What to say:

“Before I begin, I want to be explicit about what we are covering today — and what we are not. Sint Maarten’s DGTP programme has six registries. You can see all six on this slide. Today’s 40-minute session covers exactly two of them: **VROMI, the Address Registry**, and **Kadaster, the Land Registry**.

The other four — the Civil Registry, COCI Chamber of Commerce, TVSZ Tax Authority, and SZV Social Insurance — will each have their own dedicated sessions using this same methodology and dashboard format. This is not a deprioritisation; it is a sequencing decision. VROMI and Kadaster come first because address data underpins every other registry, both systems use ArcGIS, and they are tightly coupled via parcel ID linkage.

I also want to be clear: the **functional CRM** — citizen profiles, case management, service requests — is a completely separate parallel workstream. Today is exclusively the **Non-CRM integration track**.”

[PAUSE] Point to the SVG architecture diagram on the right side of the slide. The left zone (orange border) shows today’s scope. The right zone (dashed grey) shows future sessions. The ‘BLOCKED’ indicator shows current state.

If any stakeholder asks “what about the Civil Registry?” or “when do we cover COCI?” — the answer is: “Each gets its own dedicated session with the same depth and methodology. Today we need 100% of the room’s attention on VROMI and Kadaster.”

NEW: RLM Blocker Dependency Tree — Executive Summary Slide

RLM TREE — Expandable block at the bottom of Slide 2

[Click the “RECURSIVE LATENT MAPPING — BLOCKER DEPENDENCY TREE” bar to expand. The tree renders as an interactive dendrogram.]

At the bottom of the Executive Summary slide there is an expandable block titled **Recursive Latent Mapping — Blocker Dependency Tree**. This is the most technically impressive element in the presentation and particularly valuable when PhD-level or technical stakeholders are in the room.

What to say when expanding it:

“For those who appreciate a more formal view of the dependency structure, I want to show you the RLM tree. RLM — Recursive Latent Mapping — is the algorithm underlying this visualisation. It models each of our 37 blockers as a node in a directed acyclic graph, with weighted causal edges representing downstream dependencies. The algorithm performs a depth-first traversal from the integration goal root, propagating criticality weights multiplicatively down each branch.

What this produces is a precise view of *which blockers matter most* — not by opinion, but by computation. Click the plus button on any node to expand its branch and see what it blocks downstream.”

Bridge Nodes — What the B badge means

Three nodes in the tree carry an orange **B** badge. These are **bridge nodes** — blockers whose resolution unblocks the greatest number of downstream items simultaneously. They are identified by their *fan-out coefficient*: the number of nodes that become resolvable the moment this one is resolved.

Node	Item	Fan-out
§1.1	VROMI ArcGIS URL	8 items unblocked
§1.4	Kadaster ArcGIS URL	4 items unblocked
§12.2	Token endpoint URL	3 items unblocked

The remaining 34 nodes have a fan-out below 3 and carry no bridge badge. This is why §1.1 is the single highest-leverage action in the entire programme.

For PhD/technical reviewers: RLM connects formally to Critical Path Method (CPM), PERT, and fault tree analysis (FTA). Bridge node detection is analogous to articulation point identification in the spanning tree of the dependency subgraph. Acknowledging this connection impresses technical reviewers and pre-empts methodological challenge.

Fullscreen mode: Click the  FULL SCREEN button to open the tree in a dedicated full-screen overlay. In fullscreen, the **PSM Assistant**, **Reference Lookup**, and **AI Chat** are all available simultaneously alongside the tree.

NEW: Dashboard Features — Speaker Reference

This section documents the interactive features available throughout the presentation. Familiarise yourself with each before the session.

§ Reference Lookup Panel

A floating **§ References** button sits at the bottom-left of every slide. Clicking it opens a glass panel from the left containing all 43 section references from the Non-CRM Integration Readiness Checklist. Every **SX.X** tag visible in the slides and AI chat is **clickable** — clicking any reference tag opens the panel and searches for that item automatically.

The panel shows for each reference: priority badge (CRITICAL / HIGH RISK / CONFIRMED / INFO), title, description, owner, and deadline colour-coded as TODAY (red), 48 hours (amber), or This week (neutral).

Use this during Q&A: When a stakeholder asks “what exactly is §9.4?” — click the tag and the full description appears instantly without searching.

AI Chat — Ask AI Button

The dark **Ask AI** bubble at the bottom-right is available on every slide. Clicking it opens a conversation panel that knows the full integration context. Slide-specific question chips appear above the bubble automatically when you change slides — these are pre-written questions relevant to whatever slide is currently showing.

Use this during Q&A: If a question arises that you want to think through carefully, say “let me check that in the dashboard” and use the AI chat. It will give a concise, contextually accurate answer.

PSM — Persona Selection Mode (Fullscreen only)

Available when the RLM tree is in fullscreen view. Click  **PSM ASSISTANT** in the fullscreen toolbar to open a side panel. Select one of eight stakeholder personas and the AI will explain the RLM dependency tree in language tailored specifically to that role:

- **Executive / Decision Maker** — business risk, timeline, who must act
- **GoSM Legal** — Article 22, EA Ordinance, restricted fields
- **VROMI / Kadaster IT** — concrete ArcGIS admin steps, exact URLs
- **GoSM ICT** — firewall rules, TLS, IP allowlisting
- **DLT / NRPB** — programme governance, named ownership

- **PhD / Technical Reviewer** — DAG formalism, RLM methodology, graph theory
- **Programme Manager** — critical path, RAID log, sprint readiness
- **Citizen / Non-Technical** — plain English, what the portal will do

Responses are formatted as bullet points, maximum 150 words, always ending with a follow-up question. Each response includes clickable follow-up chips so stakeholders can continue without typing.

Use this in the session: If a PhD-level reviewer is present, open the PSM panel, select their persona, and let the AI explain the RLM methodology in formal graph-theoretic language while you facilitate the room discussion.

Quick Check-In Slides — Updated Format

Each of the eight engagement slides now contains **8-11 interactive questions** mixing five types: live polls (click to vote with animated result bars), drag-to-rank priority sorters, confidence sliders, word cloud tag pickers, and open text cards. Every engagement slide also contains either an SVG figure or a data table.

All questions use polite, collaborative framing. They do not feel like a knowledge test — they feel like a genuine conversation. The tone is: “I would greatly appreciate knowing...” not “Tell me the answer to...”